

Laboratory Assignment 3: Power-Performance Trade-offs

Contributors: Mehrzad Nejat, Mohammad Waqar Azhar, Per Stenstrom

Background: In the previous lab assignments, you studied the effect of different architectural parameters of a processor with a simple memory subsystem on performance. You started with a simple in-order single-issue processor with a small single-level cache and improved the performance step-by-step. First, you improved the instruction and data cache, then you increased the branch prediction accuracy, and finally you exploited ILP by allowing out-of-order execution and multiple issue in a superscalar processor. You have managed to reduce the program execution time significantly by increasing the resources. But, it doesn't come for free.

Learning outcome: In this assignment you are going to study the Power/Energy costs of performance improvements. For this purpose, you are going to design a high performance processor using the Sniper simulator [1]. Then, you will perform Power/Energy measurement using McPAT [3]. There is usually an upper limit on the power consumption of a processor called the "Power Wall". If the power goes beyond this limit, the cooling system fails to keep the chip temperature within a safe bound. So, in the next step you will review your design choices to reduce the chip power consumption below the limits. Furthermore, higher energy consumption means lower battery lifetime on portable devices and higher electricity cost for high-end computing systems. Thus, in the final step you will study the Energy-Performance trade-off in your processor model and try to improve energy efficiency.

At the end of this assignment, you should be able to:

- Measure and report the performance of a single-core processor model with sniper simulator.
- Measure and report the power/energy consumption of different components of processor and memory hierarchy using McPAT
- Understand the design choices that affect power consumption in order to meet the power limits
- Understand energy efficiency and the trade-off between energy consumption and performance

References:

- [1] <http://snipersim.org/w/The_Sniper_Multi-Core_Simulator>
- [2] <<https://groups.google.com/forum/#!forum/snipersim>>
- [3] <<http://www.hpl.hp.com/research/mcpat/>>

Evaluation: You must write a report that includes: A brief description of the problem and the method you used, important assumptions if you had, simulation results and observations, your design choices and the reason behind them. Detailed report guidelines will be available on Ping-Pong.

Simulation Environment

In this assignment, you are going to use a new simulation environment. If you are already familiar with Simple-Scalar, it should be easy for you to start working with Sniper + McPAT. You only need to use the basic sniper simulator with “--power” option that will generate power/energy results using McPAT. However, there are more advanced capabilities such as simulator APIs¹ that enable you to interact with the simulator from the program code under simulation or several tools already available in ‘tools’ and ‘scripts’ directories. But, you are not required to use these tools here. If you are interested, you can study the references for more details. The basic simulation environment is depicted in Figure 1. If you look carefully at this figure, you will notice that there are several configuration files. That is because you can pass several configurations each containing a subset of the full configuration to sniper using “-c” in the command line². It is also possible to set a single parameter from command line using “-g”. For further details check the simulator manual³. The simulator will generate a configuration file named “sim.cfg” in the simulation folder that contains the final settings of all the parameters.

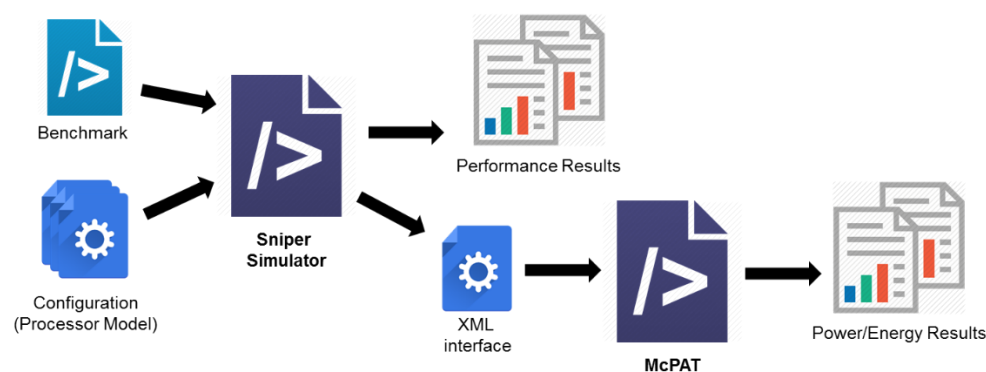


Figure 1: Simulation Environment

Similar to the previous simulations, there is a script that will run the simulator and pass the required configurations and settings. The script is named “runsim.sh” and it could be found under “Lab3” directory. You must enter the simulation home directory and the name of the configuration file (without “.cfg”) as the first and second

¹ Application Programming Interface

² If you provide multiple configurations for the same parameter, the later one will overwrite the previous ones.

³ Manuals are available in Ping-Pong documents. You can easily find them online as well.

command line arguments when executing this script. You can deactivate the simulation of each benchmark application by adding a “#” sign at the beginning of the corresponding lines in the script. You must be able to understand how this simple script works. If you have any problem understanding the parameters ask the teaching assistants and/or check the simulator manual.

Simulation Results: The performance results could be found in a file named “sim.out” while the power results are available in “power.txt” and a summary of Power/Energy measurements are printed on the standard output which is redirected to a text file¹. The energy breakdown is also plotted in “power.png”. But, you will mostly use “sim.out” and “power.txt”.

Task 1: Design a high performance processor in Sniper

You have already gained some experience and knowledge during the first two lab assignments. By now, you should be able to understand the effect of different architectural parameters of a processor and a cache on performance. In this assignment you are using a new simulator with a different simulation algorithm and modeling. The available configuration parameters are slightly different. Therefore, you cannot directly use the values from previous lab assignments. You will be provided a new base configuration file for Sniper named “base4.cfg”. One major upgrade is that now we are adding a second-level combined instruction and data cache, because most of the real processors on the market have at least two levels of cache.

1.1. Simulate the base configuration

Start by simulating the base model and collecting the performance data. Create a directory for your simulation and place the configuration file there. Run the script by passing the directory address and configuration file name as two command line arguments. Once the simulation is done, you can find the results for each benchmark application inside the corresponding sub-directory. Open one “sim.cfg” file and check the full list of configuration settings. Notice how the values of “gainestown” architecture² are overwritten by the base configuration. Also check data and tag access times to each cache. Open one “sim.out” file and check the available performance results. Fill out the first section of Table 1. Note that CPI is not directly available, but you have $IPC=1/CPI$ and also the instruction count and executed cycles. If you compare the instruction counts with the previous lab assignments, you will notice some difference. The reason is that the benchmarks were compiled differently for Simple Scalar simulations. But the instruction counts should be relatively in the same order.

¹ The gsm benchmark generates unreadable output on stdout which was redirected to a file that you never opened in the previous labs. But, now we also have Power/Energy summaries on stdout as well. So, don’t be alarmed if you see unreadable characters when you open the file that contains the redirected stdout for gsm. Actually, if you don’t see these characters, it means that something didn’t work!

² Some basic configurations such as “gainestown” are provided with the simulator in a directory named “config”

Table 1: Performance results

	Application	dijkstra	qsort	stringsearch	gsm-untoast	jpeg-cjpeg	GM*
base	Instructions						NA
	Time						NA
	CPI						NA
HPP	Time						NA
	CPI						NA
	SP**						
HPP+PW	Time						NA
	CPI						NA
	SP						

*Geometric Mean

** Speedup relative to base (use CPI values)

1.2. Design a high performance processor

Using the experience from previous assignments, modify the processor parameters to maximize the performance in a new configuration named “HPP¹”. There are a few differences compared to previous assignments. In this case, you are not modifying the branch predictor. Some other parameters such as number of functional units are hard coded in the simulator and you are not modifying them either². You have a second-level cache that can further improve the memory sub-system. Therefore, keep the level-1 cache parameters unchanged except for block size that should be the same for all the caches. The other exception is when you notice that the new block size violates the rule in Table 2 footnotes. In that case, you have to increase level-1 cache size to avoid simulation errors.

You might need to perform a few tests for designing the high performance configuration, but you are not supposed to do as much simulations as you did in the first two assignments. Check Table 2 for a list of parameters that you can modify and the possible values³. Once you have the final design, complete the corresponding sections in Table 1 and Table 2.

Hint: In order to reduce the simulation time you can deactivate the simulation of the longest benchmark when doing some of the tests. Removing “--power” argument would also save you some simulation time, but make sure that you are not going to need the power measurement results for that simulation later. You can also run the simulations in parallel to speed up. One way to do that is to create one copy of the script for each benchmark application and deactivate the rest of the applications in that script. Then run each script in a separate terminal. Another method is to execute each simulation in background.

- How does a second level cache affect the performance? Is it the same as the first level caches?

¹ High Performance Processor

² It is possible to modify the simulator since the source code is available. You can even find some suggestions and hints from the designers in the google group [2]

³ It does not necessarily mean that other values are unacceptable for the simulator. It is just a limited scope for this assignment.

Table 2: Configuration Parameters

Parameter	Possible values	HPP	HPP+PW	E-eff
in_order	True/False			
Outstanding loads and stores (LSQ)	1,8,16			
RS* entries	1,16,32			
Interval timer window size (ROB)	16,64,128			
Dispatch & Commit width	1,4,8			
L2 cache size (KB)	4,8,16,32			
L2 cache associativity	1,2,4,8			
block size** (B)	16,32,64			

* Reservation Stations

** block size should be the same for all the caches

Rule: Make sure that **for each cache** the number of sets (size / (associativity*block size)) is bigger than or equal to block size. So, you might need to increase the L1 size if you use bigger block sizes.

Task2: Power Measurements

It is time to look into the costs of the performance improvements. The two important costs are Power/Energy consumption and area overhead. Power measurements were already activated in the script by having "--power" argument. If you check a simulation output directory, you can find an image file that shows the energy break down. Detailed power values could be found in "power.txt" and "power.py" and a final Power/Energy summary was printed in the standard output and redirected to a text file. However, you are going to collect the power data from "power.txt". You can also find some estimations of area in this file. Notice the hierarchy of components. We are only interested in the core that contains the two levels of cache. So, only focus on the section related to the core. The total values could be found on top of the section followed by the detailed results for different components inside the core.

2.1. Compare the power/energy/area of base and HPP:

Look back at the simulation results for the base and HPP configurations. Collect the power and area measurement results for the core (total values) and fill out the corresponding sections of Table 3. You have noticed that there are 3 important power values¹: peak-dynamic, runtime-dynamic, and leakage. Peak-dynamic is usually used when analyzing the power wall and temperature limitations whereas runtime-dynamic is the average power consumption used for energy measurements. Leakage is another component of the power consumption that exists as long as the system is powered up, even if the system is idle with no operation. Power gating in sleep modes is one of the dynamic power management techniques used to reduce the leakage. But, we are not looking into that in this assignment.

Use the following formula to calculate the energy:

$$Energy = (P_{runtime\ dynamic} + P_{leakage}) \times Time \quad (1)$$

¹ Ignore the "power gating" results

- How big are the power, energy, and area costs of your performance improvements?
- Compare the share of different components in power, energy, and area costs.
- Is your observation the same for different benchmark applications? If not, what could be the reason?

Table 3: Power, Energy, and Area results

	Application	dijkstra	qsort	stringsearch	gsm-untoast	jpeg-cjpeg
base	Area					
	Peak Dynamic Power					
	Runtime Dynamic Power					
	Subthreshold Leakage Power					
	Energy					
HPP	Area					
	Peak Dynamic Power					
	Runtime Dynamic Power					
	Subthreshold Leakage Power					
	Energy					
	Area Increment* (%)					
	Energy Increment* (%)					
HPP+PW	Area					
	Peak Dynamic Power					
	Energy					

2.2. Meet the power wall

In this step, we introduce a power wall of **11.5 W**. Now, you probably need to make some modifications to your HPP design to make sure the “core Peak power + leakage power” remains below this limit for all the applications¹. Name your new configuration “HPP+PW” and fill Table 1, Table 2 and Table 3.

- How much did you lose performance when trying to meet the power wall?

Task3: Energy Efficiency

You have noticed a trade-off between performance and energy consumption during the last task. In this task you will further study this trade-off by using Energy-Delay-Product (EDP). EDP is one of the common metrics for measuring energy efficiency of a system. It combines both performance and energy improvement in a single parameter. Use the following formula for calculating EDP:

¹ Hint: If you manage to reduce the power for the most power hungry application, then other applications will probably meet the power wall limitation as well. So, you don’t need to simulate all the applications for intermediate tests. But you need to confirm the final result for all of them.

$$EDP = Energy \times Execution Time \quad (2)$$

First calculate the EDP for the base and HPP configurations. Then, propose one design with lower EDP and call it “E-eff”. Use your experience from the last two tasks. Fill out Table 2 and Table 4. You don’t need to consider the power wall for this task.

Table 4: Energy Efficiency Results

	Configuration	dijkstra	qsort	stringsearch	gsm-untost	jpeg-cjpeg
EDP	base					
	HPP					
	E-eff					

- Is your E-eff design closer to the base or HPP?